

# Ray Tracing and Beyond

JAMIE HONG, ANZE LIU, AKSHIT DEWAN, and TIM TU

## ACM Reference Format:

Jamie Hong, Anze Liu, Akshit Dewan, and Tim Tu. 2023. Ray Tracing and Beyond. 1, 1 (May 2023), 6 pages. <https://doi.org/10.1145/nmnnnn.nmnnnn>

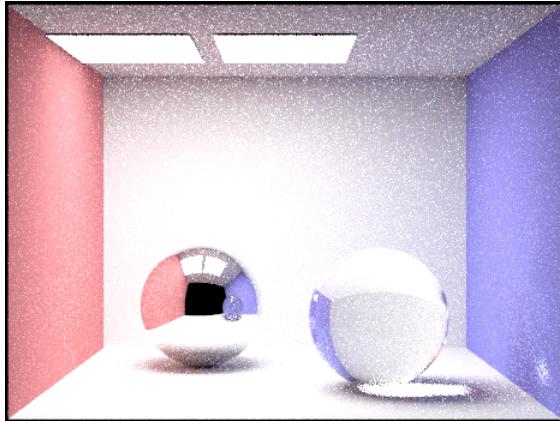


Fig. 1. Spheres with 2 lights, using photon mapping for caustics

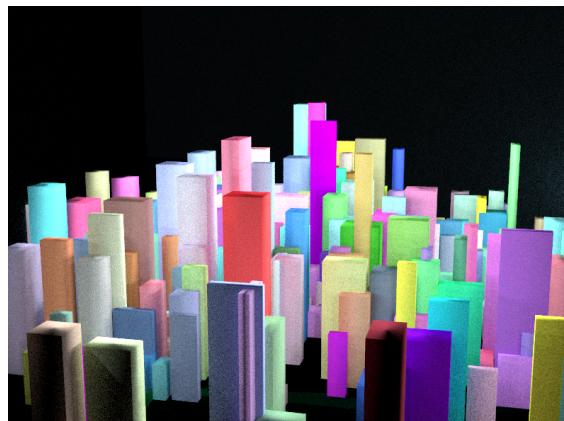


Fig. 2. Tiny city scene with 300 lights using k-d tree for cache points lookup

## 1 ABSTRACT

Ray tracing is a powerful rendering technique as it is based on simulation of the physical behavior of light. However, using ray tracing for rendering caustics for refractive objects or for scenes with thousands of lights can be inefficient.

To address these issues, we designed and implemented additional optimizations and features from scratch on top of a basic ray tracer. We chose to focus on two main additions: photon mapping, and a cache points optimization for lighting, inspired by Disney's Hyperion Renderer [1].

Along the way, this project has presented challenges in various aspects. For instance, we needed to consider how to integrate new techniques into the existing ray tracer, how to efficiently optimize our implementation, and how to best demonstrate our ray tracer's ability to render complex scenes: ones with clearer caustics or others with hundreds to thousands of lights.

With our objective and the project's challenges in mind, we augmented the basic ray tracer with photon mapping and cache point optimizations for lighting, basing our design off of the papers found in our references.

---

Authors' address: Jamie Hong; Anze Liu; Akshit Dewan; Tim Tu.

---

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from [permissions@acm.org](mailto:permissions@acm.org).

© 2023 Association for Computing Machinery.

Manuscript submitted to ACM

Manuscript submitted to ACM

## 53      2 TECHNICAL APPROACH

### 54      2.1 Photon Mapping

55  
 56      For photon mapping, we simulate photons by casting light rays into the scene, with the goal of being able to more  
 57      accurately render caustics. At ray-object intersections, we store the information about the photon at this intersection,  
 58      such as its power and direction, into a map. This occurs during a pre-rendering pass. Then, at render time, we use the  
 59      nearest few photons in our map to estimate the radiance for a shading point. For details of the implementation, we  
 60      referred to the relevant section from "A Practical Guide to Global Illumination using Photon Mapping" [4].  
 61

62      Comparison between Ray Tracing and Photon Mapping:  
 63

64      Ray Tracing	65      Photon Mapping
66      Trace rays from camera	66      Trace photons from light source
67      Rays propagate radiance	67      Photons propagate flux
68      Requires many rays and bounces for caustics	68      Dense photon map where caustics are located

70      Our two-pass photon mapping algorithm is as follows:  
 71

#### 72      Pass 1: Construction of global photon map

- 73      (1) Send rays from each light source into the scene to  
               represent the travel of photons.
- 74      (2) Check for ray-object intersections.
- 75      (3) Use Russian Roulette to decide if a photon is re-  
               flected or absorbed at point of intersection.
- 76      (4) The photon at this point of intersection is stored  
               in a k-d tree.

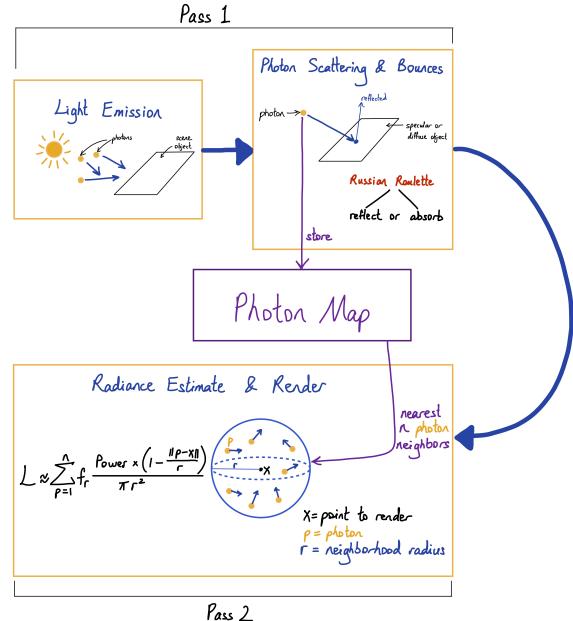
#### 77      Pass 2: Radiance estimation using photon map

- 78      (1) Based on the neighborhood radius, obtain a list  
               of photons that are neighbors to the point to be  
               rendered using k-d tree lookup.
- 79      (2) Compute radiance from each nearby photon using  
               the point's BSDF function and the photon's power,  
               scaled by the distance between the photon and  
               point.
- 80      (3) Sum the radiance estimates from the neighboring  
               photons together, as resulting radiance.
- 81      (4) Estimate radiance for every point in scene.

82      We tuned the following parameters to achieve the final rendered results.  
 83

- 84      (1) Number of rays sent into the scene to construct the photon map.
- 85      (2) Radius within which photons are used to estimate the radiance at a particular intersection point of interest.

86      The more rays we send into the scene or more photons we use to estimate radiance, the longer the scene takes to  
 87      render. Therefore, tuning the number of light rays and the neighborhood radius to achieve efficient processing while  
 88      still utilizing the photon map to render caustics is a challenge.  
 89



105 In addition, in the original photon mapping paper [4], three photon maps are constructed: caustic, global, and volume  
 106 photon maps. In our project, we implemented only the global photon map to achieve an approximate representation of  
 107 the global illumination for the scene while still keeping the render time reasonable.  
 108

## 110 2.2 Cache Points Optimization

111 In path tracing with importance sampling, we loop over every light source to compute the radiance estimate for each  
 112 point. However, considering every light becomes more expensive as the number of light in the scene increases, especially  
 113 if some of the lights far away from the shading point have negligible contribution or are occluded.  
 114

115 In order to render scenes with hundreds to thousands of lights, we would like to perform efficient light sampling. One  
 116 method of doing so is to "cache" region-specific lighting information into a certain number of cache points. From these,  
 117 we can then use the nearest cache point to sample light for a shading point, resulting in more efficient computation  
 118 during the rendering phase.  
 119

120 As for the basis of this idea, "The Design and Evolution of Disney's Hyperion Renderer" [1] introduces cache points  
 121 optimization for their goal of rendering scenes with hundreds of thousands of lights, such as in ones in the movie *Big*  
 122 *Hero 6*. We also referred to Pixar's "Production Volume Rendering" [2], for additional subtleties in the design.  
 123

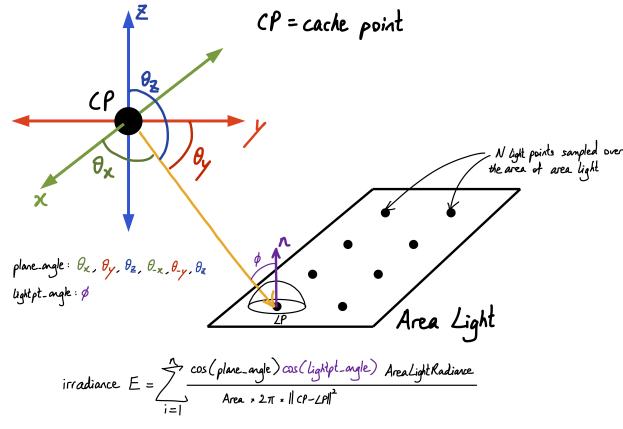
124 Our final design consisted of a 2-stage cache point implementation:

### 125 Stage 1: Pre-compute phase: building cache points

- 126 (1) Initialize uniformly randomly distributed cache  
 127 points within bounding boxes.
- 128 (2) For each cache point, create 8 light lists, each repre-  
 129 sents a discrete distribution over light sources,  
 130 based on their incident irradiance and visibility  
 131 (calculated by sampling rays from lights).
  - 132 (a) 6 lists of lights and irradiances for the cardinal  
 133 planes (1 list for each directional plane).
  - 134 (b) 1 list of lights and irradiances for the omnidi-  
 135 rectional receiver.
  - 136 (c) 1 list of lights which are very close to the  
 137 cache point, to account for rapid changes in  
 138 illumination between cache points.
- 139 (3) Filter out insignificant lights by only storing the  
 140 most significant lights until 97% of total irradiance  
 141 at this cache point is reached.
- 142 (4) The cache points are stored in a k-d tree.

### 143 Stage 2: Light sampling and aggregating phase: using cache points

- 144 (1) For each point of interest, sample light from the closest cache point using the k-d tree, which returns the nearest  
 145 neighbor [3].
- 146 (2) Randomly choose a distribution from cardinal, omnidirectional, and nearby light distributions, by weighting  
 147 each distribution by the dot product of its normal (when applicable) with the surface normal of shading point.
- 148 (3) Randomly sample a light from the chosen discrete distribution, weighting by the visibility of the light.



Incident irradiance from area light to cache point.

157 (4) Use the chosen light to generate more samples for estimating radiance at a point of intersection. Sampling from  
 158 lights in this way allows us to use cache points for importance sampling.  
 159

160 Parameters tuned to achieve quality render while lowering computation time:

- 161 (1) Number of cache points (per bounding box)  
 162 (2) Number of lights used for cache sampling  
 163 (3) Number of samples per pixel

164 As the number of cache points, lights, and samples per pixel increases, the quality of the render improves. The higher  
 165 the density of cache points, the more accurately the nearest cache point can approximate the relevant lighting at a  
 166 shading point. However, improved accuracy comes at the cost of increased render time for a scene, as lookup time could  
 167 increase with more cache points.  
 168

### 169 2.3 Challenges and Resolutions

170 2.3.1 *Scenes and Procedural Generation.* In the process of implementing our optimizations, we realized that the original  
 171 scenes from the project were likely insufficient to demonstrate the results of our optimizations. We approached this in  
 172 two main ways.  
 173

174 One method we took was to start from the existing .dae files included in the project. We updated the existing scene  
 175 to change the setup of the lights from 1 large area light, to 2 lights, and then to 440 smaller lights.  
 176

177 We were also interested in producing our own scenes, and looked into generating Blender scenes from scratch using  
 178 Blender's Python interface. Once we had written a script for procedurally generating city scenes, with many buildings  
 179 (cubes) and many lights (area lights), we could then use the script to generate various .dae files that we could then use  
 180 to render effects such as caustics or demonstrate the efficiency of our renderer with a large amount of lights.  
 181

182 We created approximately 5 city scenes, with the number of lights ranging from 300 to 5000 lights. With a larger  
 183 number of lights in each scene, we were able to better observe improvements in render time for our cache point  
 184 optimized renderer in comparison to the original basic ray tracer implementation.  
 185

186 2.3.2 *Data Structure for Cache Points Lookup.* We experimented with various storage methods for efficient neighbor  
 187 lookup of cache points. After profiling the rendering times and comparing the resulting renders of the different methods,  
 188 we decided to use a k-d tree, which enabled shorter render times with minimal artifacts.  
 189

190 Vector	191 K-D Tree	192 Spatial Map
191 Expensive sort	192 Optimize cache point lookup with few artifacts	193 Checkered pattern artifacts (10)

### 194 2.4 Lessons Learned

195 Over the course of this project, we learned how to interpret and transfer methods described in research papers to fully  
 196 functional implementations, under the constraints of limited resources and time.  
 197

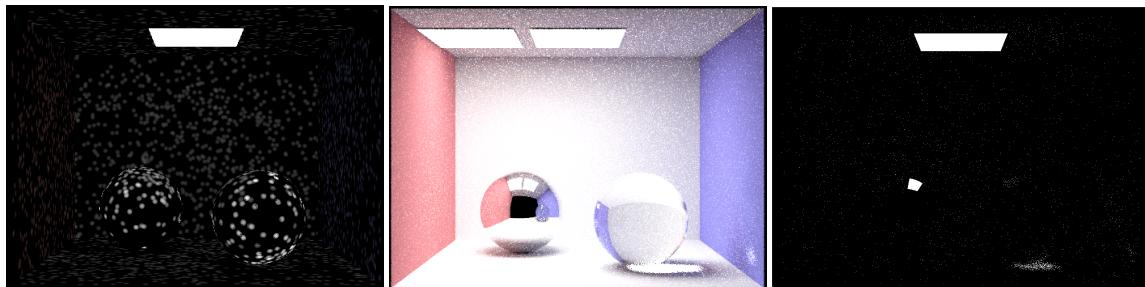
198 For example, we extracted the key steps of the photon mapping algorithm: constructing the photon map, Russian  
 199 Roulette for photon scattering, and radiance estimation. We determined that cone filtering, Gaussian filtering, and  
 200 caustic and volume photon maps would be useful for further improving render results, but were not key for our  
 201 fundamental photon mapping implementation.  
 202

203 We utilized Blender and its Python API to procedurally generate city models with buildings of refractive or reflective  
 204 material, and illuminated by area lights randomly dispersed throughout the scene, for demonstrating our cache point  
 205 optimization results.  
 206

207 Finally, we learned how different parameters can introduce tradeoffs between render time and accuracy, and how to  
 208 optimize our program to balance reducing computation time while maintaining high-quality renders.

209 **3 RESULTS**

210 **3.1 Rendered Scenes with Photon Mapping**

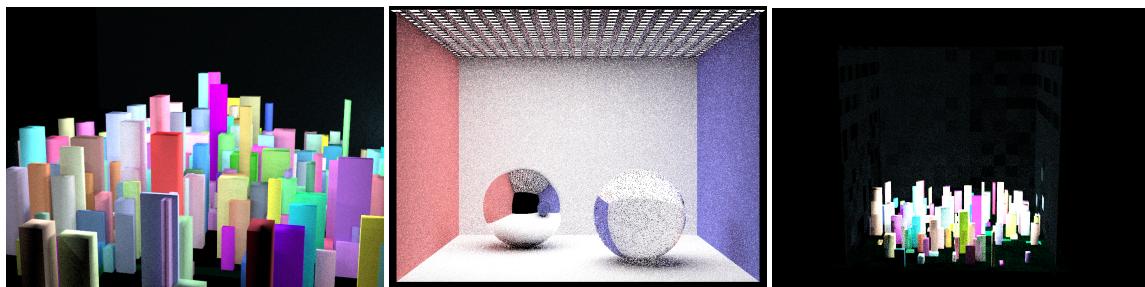


222 Fig. 3. Spheres - photon mapping with one bounce radiance

223 Fig. 4. Spheres - photon mapping with six bounce radiance and cache point optimization

224 Fig. 5. Spheres - photon map

226 **3.2 Rendered Scenes with Cache Points**

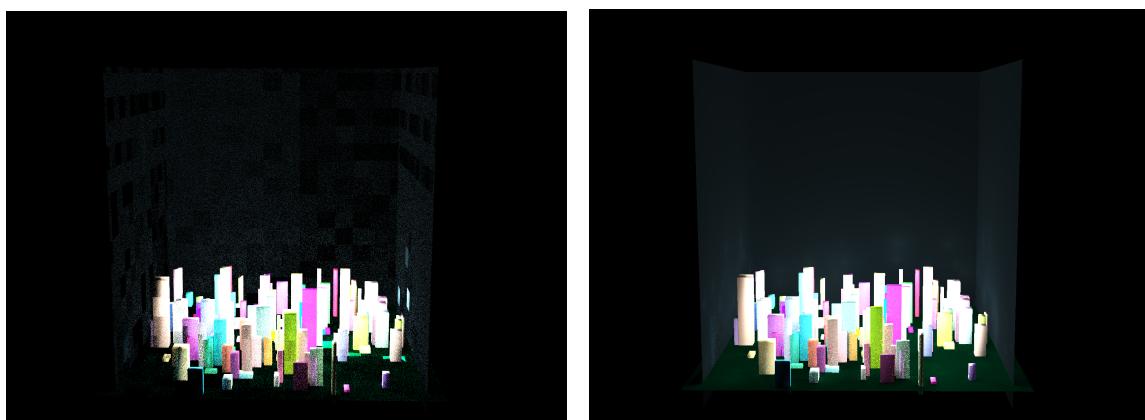


238 Fig. 6. Tiny city scene with 300 lights using kd tree for cache points lookup

239 Fig. 7. Spheres with 440 lights using k-d tree for cache points lookup

240 Fig. 8. Tiny city scene with 5000 lights using spatial map for cache points lookup

241 **3.3 Speedup using Cache Point Optimization**



258 Fig. 9. Tiny city scene with 5000 lights using spatial map for cache points lookup (render time: 115.59 s)

259 Fig. 10. Tiny city scene with 5000 lights using staff renderer for comparison (render time: 375.49 s)

**261      4 CONCLUSION**

262  
 263 Efficiently rendering complex scenes, such as those with a lot of caustics and light sources, requires complex optimizations.  
 264 We have explored and implemented two such optimizations: photon mapping and cache points. Yet ultimately,  
 265 we are only making approximations of the physical behavior of lighting through these methods. As a result, our  
 266 optimizations may trade off image quality or low amounts of noise to achieve faster rendering times. From this project  
 267 onwards, future optimizations can be further explored to expand our renderer's ability to that of production level  
 268 renderers such as Disney's Hyperion, and to allow us to render even more exciting scenes than before!  
 269

**270      5 CONTRIBUTIONS FROM EACH TEAM MEMBER**

- 271  
 272 (1) Jamie: pseudocode and implementation for cache point and photon mapping, structure used to store cache  
 273 points (kd-tree), Blender script for procedural scene generation, tuned and tested with different parameters, final  
 274 video, final report  
 275  
 276 (2) Tim: implementation of photon mapping, cache point construction, structure used to store cache points (spatial  
 277 map), light sampling, generated .dae files to be rendered, final video  
 278  
 279 (3) Akshit: implementation of photon mapping construction, cache point light sampling, generated .dae files to be  
 280 rendered, tuned and tested with different parameters, final video  
 281  
 282 (4) Anze: pseudocode for cache point and photon mapping, implementation for cache points construction and light  
 283 sampling, milestone report, final video, slides, final report

**284      REFERENCES**

- 285  
 286 [1] Brent Burley, David Adler, Matt Jen-Yuan Chiang, Hank Driskill, Ralf Habel, Patrick Kelly, Peter Kutz, Yining Karl Li, and Daniel Teece. 2018. The  
 287 Design and Evolution of Disney's Hyperion Renderer. *ACM Trans.Graph.* 37, 3 (July 2018). <https://doi.org/10.1145/3182159>  
 288 [2] Julian Fong, Magnus Wrenninge, Christopher Kulla, and Ralf Habel. 2017. Production Volume Rendering: SIGGRAPH 2017 Course. In *ACM  
 289 SIGGRAPH 2017 Courses* (Los Angeles, California) (*SIGGRAPH '17*). Association for Computing Machinery, New York, NY, USA, Article 2, 79 pages.  
 290 <https://doi.org/10.1145/3084873.3084907>  
 291 [3] J. Frederico Carvalho 2008. kd-tree implementation. <https://github.com/crvs/KDTree>.  
 292 [4] Henrik Wann Jensen. 2002. A Practical Guide to Global Illumination using Photon Mapping. *Siggraph 2002 Course 43* (July 2002).